

Jianshen Liu

SOFTWARE ENGINEER · COMPUTATIONAL STORAGE, SMARTNIC, AND DISTRIBUTED SYSTEM EXPERT

Cupertino, CA 95014, USA

☎ +1 (415) 728-6512 | ✉ ljishen@gmail.com | 🏠 www.jianshenliu.com | 📞 ljishen | 🌐 jianshenliu

Skills

DevOps	AWS/GCP, Jupyter Notebook, Docker, Kubernetes, Vagrant, Jenkins, Git, Mercurial, Maven, Gradle, VSCode, IntelliJ IDEA, CI/CD
Back-end	Linux Kernel, eBPF, DPDK, Ceph, Hadoop, MySQL, PostgreSQL, RocksDB, Apache Arrow, Ansible, REST API, NVMe, Microservices, Machine Learning, Distributed Systems
Programming	C++, Python, Java, Bash, SQL, PHP, C#, Javascript, LaTeX, Vimscript
Languages	English, Cantonese, Mandarin
Others	Patience, Open-Mindedness, Active Listening, Versatility, Enthusiasm, Agile Methodologies, Best Practices

Work Experience

University of California, Santa Cruz

Santa Cruz, USA

GRADUATE STUDENT RESEARCHER

Jan 2016 - Present

- Developing a **SmartNICs-based distributed data-flow management platform** for HPC big data workloads at Sandia National Laboratories
- Implemented *Bitar*, a **C++** library based on **Apache Arrow** and **DPDK**, to accelerate (>8x) data compression with hardware using the **BlueField-2 DPU**
- Characterized the potential of data processing using SmartNICs with microbenchmarks, **DPDK**, and an **in-kernel packet generator**; work highlighted by The Next Platform and Data Centre Dynamics
- Contributed features and bug fixes to multiple open-source projects, including **Ceph** and **Apache Arrow**
- Proposed a novel **media-based methodology** to quantify the benefits of offloading data access functions to storage devices, prototyped with **YCSB** and **RocksDB** on **ROCKPro64**, and automated with **Ansible** and **Python**
- Developed *no_fscache*, a **Linux kernel module** to address page cache pollution before RWF_UNCACHED was introduced into the POSIX API

KIOXIA America Inc.

San Jose, USA

STRATEGIC MARKETING ENGINEER INTERN

Jun 2020 - Sep 2020

- Quantified the performance speedup (>10x) of offloading SQL query processing to SSDs using **MariaDB**, **Docker**, **eBPF**, **FlameGraph**, **Linux perf**, **cgroups**, and the **TPC-H benchmark**
- Reported monthly to the senior vice president to drive internal strategic planning for **computational storage device** products
- Collaborated on cross-team projects to identify IO performance issues in NUMA systems

Samsung Semiconductor

San Jose, USA

PERFORMANCE ANALYSIS ENGINEER INTERN

Jul 2019 - Sep 2019

- Established a **distributed testing framework** to automate the performance evaluation for Samsung Network KV devices
- Created **eBPF**- and **kprobe**-based tools to expose the **performance observability** of the device API
- Identified major system bottlenecks with the tools developed, resolving issues boosted the write performance by 60%

Wumii Technology

Shenzhen, China

ANDROID DEVELOPMENT TEAM LEAD

Dec 2012 - Jun 2014

- Led design and code reviews of the first **reading recommendation app** (among the top 25 social apps) in China
- Implemented an energy-efficient and highly available Paho **MQTT**-based notification module in **Java** that outperformed commercial solutions in terms of latency by 35%

SENIOR SOFTWARE ENGINEER

Jun 2010 - Dec 2012

- Worked as a core engineer to develop the first **news recommendation system** in China
- Drove the system architecture to handle distributed requests based on the **Spring** and **Hibernate** framework, and the code obfuscation system using Google's **Closure Compiler**
- Responsible for maintaining high availability for all backend services
- Launched a **Java**-, **PHP**-, and **Javascript**-based platform to present picturized article recommendations across multiple platforms, achieved 200 million monthly active users

Education

University of California, Santa Cruz

Santa Cruz, USA

DOCTOR OF PHILOSOPHY (PH.D.), COMPUTER SCIENCE AND ENGINEERING

2016 - Present

PhD Candidate advised by Prof. Carlos Maltzahn with a GPA of 3.96/4.00

Dissertation topic: **Eusocial Storage Devices**

Honors & Awards

Sep 2022	Outstanding Student Paper , IEEE High Performance Extreme Computing (HPEC)	Virtual Conference
Oct 2009	National Encouragement Scholarship , South China Agricultural University	Guangzhou, China
Oct 2009	The Second Prize Scholarship , South China Agricultural University (top 10%)	Guangzhou, China
Oct 2008	State Grants , South China Agricultural University	Guangzhou, China
Oct 2008	The Second Prize Scholarship , South China Agricultural University (top 10%)	Guangzhou, China

Publications

Processing Particle Data Flows with SmartNICs

Jianshen Liu, Carlos Maltzahn, Matthew Leon Curry, Craig Ulmer

26th Annual 2022 IEEE High Performance Extreme Computing (IEEE-HPEC 2022), 2022, Virtual Conference

Performance Characteristics of the BlueField-2 SmartNIC

Jianshen Liu, Carlos Maltzahn, Craig Ulmer, Matthew Leon Curry

Sandia National Lab. (SNL-CA), Livermore, CA (United States); Sandia National Lab. (SNL-NM), Albuquerque, NM (United States) (May 2021). 2021

Implementing a Kernel Module to Eliminating External Caching Effects

Jianshen Liu

The Center for Research in Open Source Software (CROSS) (Jan. 2020). 2020

Scale-out Edge Storage Systems with Embedded Storage Nodes to Get Better Availability and {Cost-Efficiency} At the Same Time

Jianshen Liu, Matthew Leon Curry, Carlos Maltzahn, Philip Kufeldt

3rd USENIX Workshop on Hot Topics in Edge Computing (HotEdge'20), 2020, Virtual Conference

MBWU: Benefit Quantification for Data Access Function Offloading

Jianshen Liu, Philip Kufeldt, Carlos Maltzahn

International Conference on High Performance Computing (HPC-IODC 2019), 2019, Springer, Cham

Presentations & Posters

Experience of using the BlueField-2 DPU

Presented at Sandia National Laboratories (May 2022). 2022

Processing Particle Data Flows with SmartNICs

Presented at UC Santa Cruz Open Source Symposium (Sept. 2022). 2022

Eusocial Storage Devices

Presented at KIOXIA America Inc., Seagate Technology, and Fujitsu America Inc. (2020). 2020

MBWU (MibeeWu): Quantifying benefits of offloading data management to storage devices

Poster Session at 17th USENIX Conference on File and Storage Technologies (FAST'19) (Feb. 2019). 2019

Quantifying benefits of offloading data management to storage devices

Poster Session at 2019 OCP Global Summit Symposium (Mar. 2019). 2019

Certificates

Jan 2022	CruzHacks Mentor Certificate , CruzHacks	Santa Cruz, USA
Dec 2020	Graduate Student Professional Communication Certificate , University of California, Santa Cruz	Santa Cruz, USA

Program Committees

CruzHacks

TECH MENTOR

Mentored hackers to build web applications to solve real-world problems

Santa Cruz, USA

Jan 2020, Jan 2021

Center for Research in Open Source Software (CROSS)

WORKSHOP CHAIR

Chaired the SmartNICs and Eusocial Storage Devices workshop at the 5th and 6th CROSS Research Symposium at UC Santa Cruz

Santa Cruz, USA

Oct 2021, Oct 2020